

In the name of God.

```
In [1]: # Step 1: Install required libraries
!pip install -q kaggle
!pip install -q opencv-python
!pip install -q numpy

# Step 2: Upload the kaggle.json file
from google.colab import files
files.upload() # Upload the kaggle.json file here

# Step 3: Create folder for kaggle and transfer file
import os
os.makedirs('/root/.kaggle', exist_ok=True)
!cp kaggle.json /root/.kaggle/

# Step 4: Set permissions
!chmod 600 /root/.kaggle/kaggle.json

# Step 5: Download the dataset from Kaggle
!kaggle datasets download -d paultimothymooney/chest-xray-pneumonia

# Step 6: Extract the ZIP file
import zipfile

with zipfile.ZipFile('chest-xray-pneumonia.zip', 'r') as zip_ref:
    zip_ref.extractall('/content/chest-xray-pneumonia')

# Step 7: Check the directory structure
import os

# Display available directories
print("Directory structure after extraction:")
for root, dirs, files in os.walk('/content/chest-xray-pneumonia'):
    print(f"Folder: {root} - Files: {files}")
```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

```
Saving kaggle (1).json to kaggle (1).json
cp: cannot stat 'kaggle.json': No such file or directory
chmod: cannot access '/root/.kaggle/kaggle.json': No such file or directory
Dataset URL: https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia (https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia)
License(s): other
Downloading chest-xray-pneumonia.zip to /content
100% 2.29G/2.29G [00:10<00:00, 228MB/s]
100% 2.29G/2.29G [00:10<00:00, 224MB/s]
Directory structure after extraction:
Folder: /content/chest-xray-pneumonia - Files: []
Folder: /content/chest-xray-pneumonia/chest_xray - Files: []
Folder: /content/chest-xray-pneumonia/chest_xray/val - Files: []
```

```
In [2]: import shutil
```

```
# Remove extra folders
shutil.rmtree('/content/chest-xray-pneumonia/chest_xray/__MACOSX', ignore_errors=True)
shutil.rmtree('/content/chest-xray-pneumonia/chest_xray/chest_xray', ignore_errors=True)

print("Deleted extra folders.")
```

```
Deleted extra folders.
```

In [3]:

```
import os
import cv2
import numpy as np

def load_images(data_path):
    images = []
    labels = []

    for label in ['NORMAL', 'PNEUMONIA']:
        label_path = os.path.join(data_path, label)
        if not os.path.exists(label_path):
            print(f'The file does not exist: {label_path}')
            continue

        for img_file in os.listdir(label_path):
            img_path = os.path.join(label_path, img_file)
            img = cv2.imread(img_path)
            if img is None:
                print(f'Image failed to load: {img_path}')
                continue
            img = cv2.resize(img, (128, 128)) # Resize to 128x128
            images.append(img)
            labels.append(int(label == 'PNEUMONIA')) # Labels are numerical

    return np.array(images), np.array(labels)

# Upload images
data_path = '/content/chest-xray-pneumonia/chest_xray'
images, labels = load_images(data_path)

print(f'Number of images uploaded: {len(images)}')
print(f'Dimensions of each image: {images.shape[1:]}')
print(f'Tags: {labels[:10]}') # Showing the first 10 tags
```

The file does not exist: /content/chest-xray-pneumonia/chest_xray/NORMAL
The file does not exist: /content/chest-xray-pneumonia/chest_xray/PNEUMONIA
Number of images uploaded: 0
Dimensions of each image: ()
Tags: []

In [4]: `import os`

```
# Show directory structure
base_path = '/content/chest-xray-pneumonia'
for dirname, dirnames, filenames in os.walk(base_path):
    print(f'Folder: {dirname}, Files: {filenames}')


Folder: /content/chest-xray-pneumonia, Files: []
Folder: /content/chest-xray-pneumonia/chest_xray, Files: []
Folder: /content/chest-xray-pneumonia/chest_xray/val, Files: []
Folder: /content/chest-xray-pneumonia/chest_xray/val/PNEUMONIA, Files: ['person1951_bacteria_4882.jpeg', 'person1952_bacteria_4883.jpeg', 'person1946_bacteria_4874.jpeg', 'person1950_bacteria_4881.jpeg', 'person1946_bacteria_4875.jpeg', 'person1947_bacteria_4876.jpeg', 'person1954_bacteria_4886.jpeg', 'person1949_bacteria_4880.jpeg']
Folder: /content/chest-xray-pneumonia/chest_xray/val/NORMAL, Files: ['NORMAL2-IM-1442-0001.jpeg', 'NORMAL2-IM-1437-0001.jpeg', 'NORMAL2-IM-1440-0001.jpeg', 'NORMAL2-IM-1427-0001.jpeg', 'NORMAL2-IM-1436-0001.jpeg', 'NORMAL2-IM-1430-0001.jpeg', 'NORMAL2-IM-1431-0001.jpeg', 'NORMAL2-IM-1438-0001.jpeg']
Folder: /content/chest-xray-pneumonia/chest_xray/test, Files: []
Folder: /content/chest-xray-pneumonia/chest_xray/test/PNEUMONIA, Files: ['person150_bacteria_715.jpeg', 'person122_bacteria_583.jpeg', 'person3_virus_16.jpeg', 'person175_bacteria_833.jpeg', 'person46_virus_96.jpeg', 'person59_virus_116.jpeg', 'person120_bacteria_570.jpeg', 'person1633_virus_2829.jpeg', 'person1608_virus_2786.jpeg', 'person120_bacteria_572.jpeg', 'person117_bacteria_553.jpeg', 'person155_bacteria_730.jpeg', 'person85_bacteria_424.jpeg', 'person96_bacteria_466.jpeg', 'person85_bacteria_419.jpeg', 'person1673_virus_2889.jpeg', 'person161_bacteria_759.jpeg', 'person117_bacteria_557.jpeg', 'person100_bacteria_478.jpeg', 'person21_bacteria_579.jpeg', 'person1647_virus_2848.jpeg', 'person71_virus_132.jpeg', 'person143_bacteria_689.jpeg', 'person22_virus_55.jpeg', 'person96_bacteria_465.jpeg', 'person80_bacteria_392.jpeg', 'person136_bacteri
```

In [5]: `import os`

```
data_path = '/content/chest-xray-pneumonia/chest_xray'
categories = ['train', 'val', 'test']
subfolders = ['NORMAL', 'PNEUMONIA']

for category in categories:
    for subfolder in subfolders:
        folder_path = os.path.join(data_path, category, subfolder)
        if os.path.exists(folder_path):
            files = os.listdir(folder_path)
            print(f'Folder: {folder_path} - Number of files: {len(files)}')
        else:
            print(f'The folder does not exist: {folder_path}')
```

Folder: /content/chest-xray-pneumonia/chest_xray/train/NORMAL - Number of files: 1341
Folder: /content/chest-xray-pneumonia/chest_xray/train/PNEUMONIA - Number of files: 3875
Folder: /content/chest-xray-pneumonia/chest_xray/val/NORMAL - Number of files: 8
Folder: /content/chest-xray-pneumonia/chest_xray/val/PNEUMONIA - Number of files: 8
Folder: /content/chest-xray-pneumonia/chest_xray/test/NORMAL - Number of files: 234
Folder: /content/chest-xray-pneumonia/chest_xray/test/PNEUMONIA - Number of files: 390

```
In [6]: import os
import cv2
import numpy as np

def load_images(data_path):
    images = []
    labels = []

    for category in ['train', 'val', 'test']:
        for label in ['NORMAL', 'PNEUMONIA']:
            folder_path = os.path.join(data_path, category, label)
            if os.path.exists(folder_path):
                for img_file in os.listdir(folder_path):
                    img_path = os.path.join(folder_path, img_file)
                    img = cv2.imread(img_path)
                    if img is not None:
                        img = cv2.resize(img, (128, 128)) # Resize to 128x128
                        images.append(img)
                        labels.append(0 if label == 'NORMAL' else 1) # Labels: 0 for normal and 1 for pneumonia
                    else:
                        print(f'Image failed to load: {img_path}')
            else:
                print(f'The folder does not exist: {folder_path}')

    return np.array(images), np.array(labels)

# Upload images
data_path = '/content/chest-xray-pneumonia/chest_xray'
images, labels = load_images(data_path)

print(f'Number of images uploaded: {len(images)}')
print(f'Dimensions of each image: {images[0].shape}' if len(images) > 0 else 'Dimensions of each image: (no images)')
print(f'Tags: {np.unique(labels, return_counts=True)}')
```

```
Number of images uploaded: 5856
Dimensions of each image: (128, 128, 3)
Tags: (array([0, 1]), array([1583, 4273]))
```



```
In [7]: from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical

# Dividing data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(images, labels, test_size=0.2, random_state=42)

# Converting tags to category format (one-hot encoding)
y_train = to_categorical(y_train, 2)
y_test = to_categorical(y_test, 2)

# Construction of CNN model
model = Sequential()

# Convolutional layers and pooling
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

# Dense and Dropout layers
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(2, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Model training
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))

# Model evaluation
loss, accuracy = model.evaluate(X_test, y_test)
```

```
print(f'Model Accuracy: {accuracy * 100:.2f}%')
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/10

```
147/147 ━━━━━━━━━━ 22s 63ms/step - accuracy: 0.7199 - loss: 26.8553 - val_accuracy: 0.8882 - val_loss: 0.2735
```

Epoch 2/10

```
147/147 ━━━━━━━━ 3s 17ms/step - accuracy: 0.8829 - loss: 0.3045 - val_accuracy: 0.9019 - val_loss: 0.2489
```

Epoch 3/10

```
147/147 ━━━━━━ 5s 16ms/step - accuracy: 0.9212 - loss: 0.2319 - val_accuracy: 0.9206 - val_loss: 0.2017
```

Epoch 4/10

```
147/147 ━━━━ 2s 15ms/step - accuracy: 0.9420 - loss: 0.1587 - val_accuracy: 0.9283 - val_loss: 0.2045
```

Epoch 5/10

```
147/147 ━━━━ 2s 15ms/step - accuracy: 0.9445 - loss: 0.1547 - val_accuracy: 0.9480 - val_loss: 0.1474
```

Epoch 6/10

```
147/147 ━━━━ 3s 17ms/step - accuracy: 0.9570 - loss: 0.1196 - val_accuracy: 0.9488 - val_loss: 0.1571
```

Epoch 7/10

```
147/147 ━━━━ 5s 16ms/step - accuracy: 0.9620 - loss: 0.1071 - val_accuracy: 0.9386 - val_loss: 0.1750
```

Epoch 8/10

```
147/147 ━━━━ 2s 15ms/step - accuracy: 0.9661 - loss: 0.0956 - val_accuracy: 0.9386 - val_loss: 0.1714
```

Epoch 9/10

```
147/147 ━━━━ 3s 15ms/step - accuracy: 0.9640 - loss: 0.0926 - val_accuracy: 0.9505 - val_loss: 0.1752
```

Epoch 10/10

```
147/147 ━━━━ 3s 17ms/step - accuracy: 0.9661 - loss: 0.0810 - val_accuracy: 0.9437 - val_loss: 0.1882
```

```
37/37 ━━━━ 0s 5ms/step - accuracy: 0.9395 - loss: 0.1975
```

Model Accuracy: 94.37%

In [8]: # Precision vs Recall

```
import matplotlib.pyplot as plt
from sklearn.metrics import precision_recall_curve, classification_report, confusion_matrix
import numpy as np

# model training
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))

# Predictions
y_pred = model.predict(X_test)

# Convert one-hot encoded labels to original labels
y_test_labels = np.argmax(y_test, axis=1)
y_pred_labels = np.argmax(y_pred, axis=1)

# Calculation of precision and recall
precision, recall, thresholds = precision_recall_curve(y_test_labels, y_pred_labels)

# Full report including precision, recall, f1-score and accuracy
print(classification_report(y_test_labels, y_pred_labels))

# Draw a precision vs recall graph
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, color='magenta')
plt.title('Precision vs Recall', fontsize=14)
plt.xlabel('Recall', fontsize=12)
plt.ylabel('Precision', fontsize=12)
plt.grid(True)
plt.show()

# Calculating the clutter matrix to see the number of True Positives, True Negatives, False Positives, False Neg
conf_matrix = confusion_matrix(y_test_labels, y_pred_labels)
print("Confusion Matrix:")
print(conf_matrix)
```

Epoch 1/10

147/147  2s 16ms/step - accuracy: 0.9713 - loss: 0.0815 - val_accuracy: 0.9471 - val_loss: 0.1863

Epoch 2/10

147/147  2s 15ms/step - accuracy: 0.9764 - loss: 0.0599 - val_accuracy: 0.9462 - val_loss: 0.2145

Epoch 3/10

147/147  3s 15ms/step - accuracy: 0.9706 - loss: 0.0760 - val_accuracy: 0.9445 - val_loss: 0.2042

Epoch 4/10

147/147  3s 16ms/step - accuracy: 0.9791 - loss: 0.0595 - val_accuracy: 0.9514 - val_loss: 0.1638

Epoch 5/10

147/147  2s 16ms/step - accuracy: 0.9853 - loss: 0.0444 - val_accuracy: 0.9317 - val_loss: 0.2396

Epoch 6/10

147/147  2s 15ms/step - accuracy: 0.9622 - loss: 0.1045 - val_accuracy: 0.9514 - val_loss: 0.2055

Epoch 7/10

147/147  3s 16ms/step - accuracy: 0.9743 - loss: 0.0694 - val_accuracy: 0.9420 - val_loss: 0.2358

Epoch 8/10

147/147  2s 16ms/step - accuracy: 0.9785 - loss: 0.0622 - val_accuracy: 0.9522 - val_loss: 0.1909

Epoch 9/10

147/147  3s 16ms/step - accuracy: 0.9843 - loss: 0.0445 - val_accuracy: 0.9428 - val_loss: 0.2292

Epoch 10/10

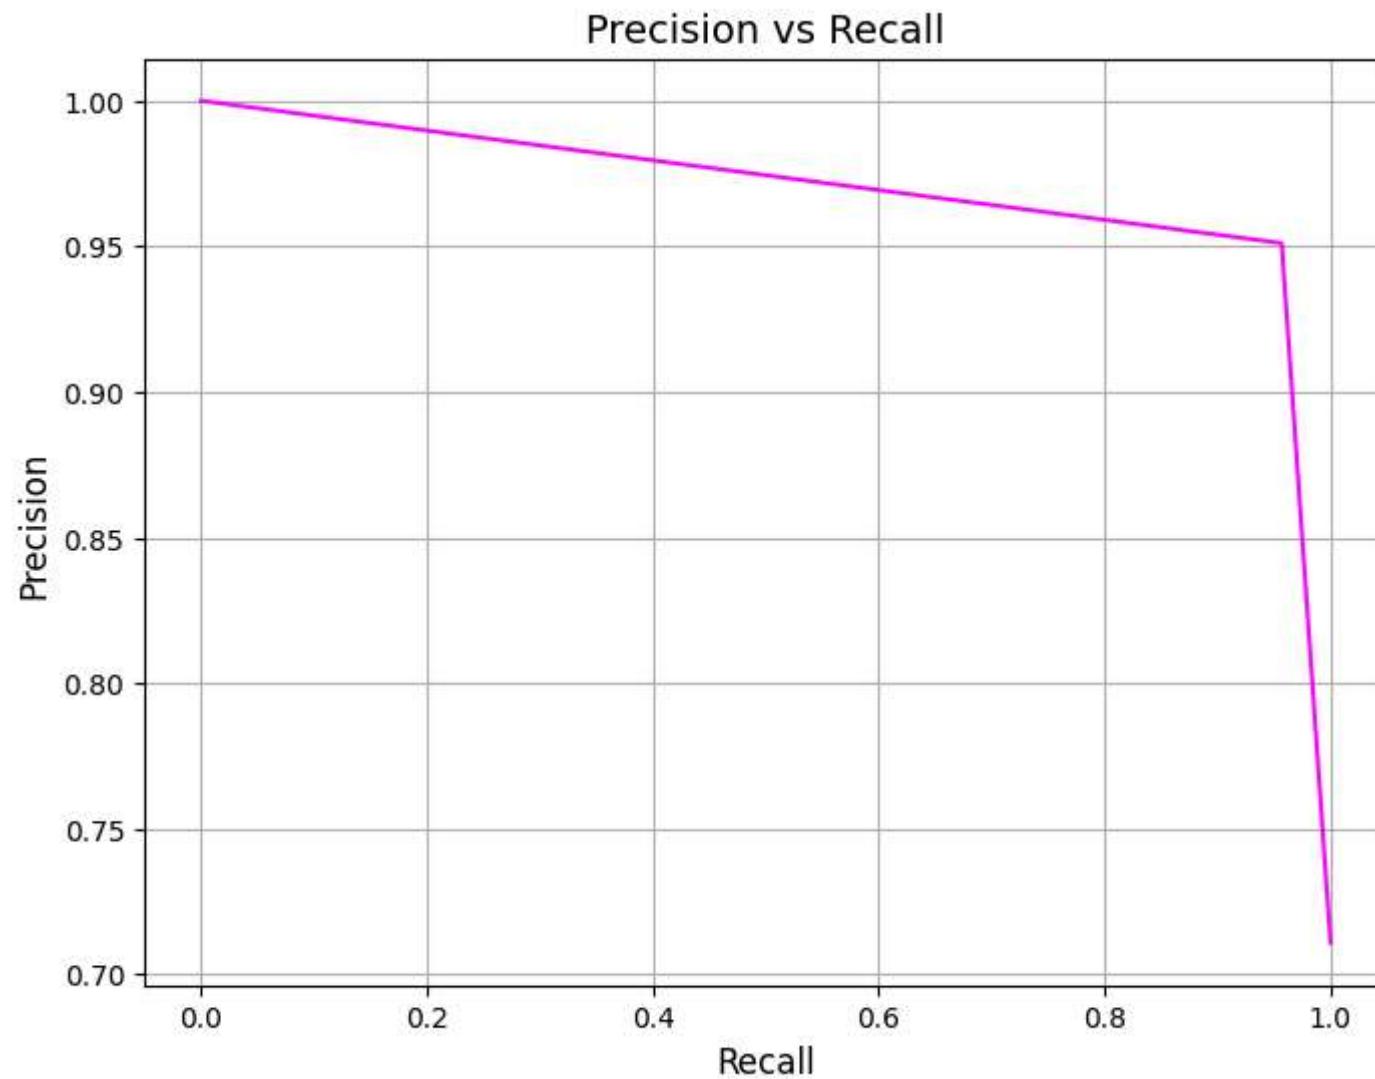
147/147  3s 16ms/step - accuracy: 0.9823 - loss: 0.0437 - val_accuracy: 0.9343 - val_loss: 0.2508

37/37  1s 10ms/step

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.89	0.88	0.89	339
1	0.95	0.96	0.95	833

accuracy			0.93	1172
macro avg	0.92	0.92	0.92	1172
weighted avg	0.93	0.93	0.93	1172



Confusion Matrix:

```
[[298  41]
 [ 36 797]]
```

In [9]: # Save the model

```
model.save('pneumonia_detection_model.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
In [10]: import matplotlib.pyplot as plt

# Suppose history contains the history of model training
# history = model.fit(...)

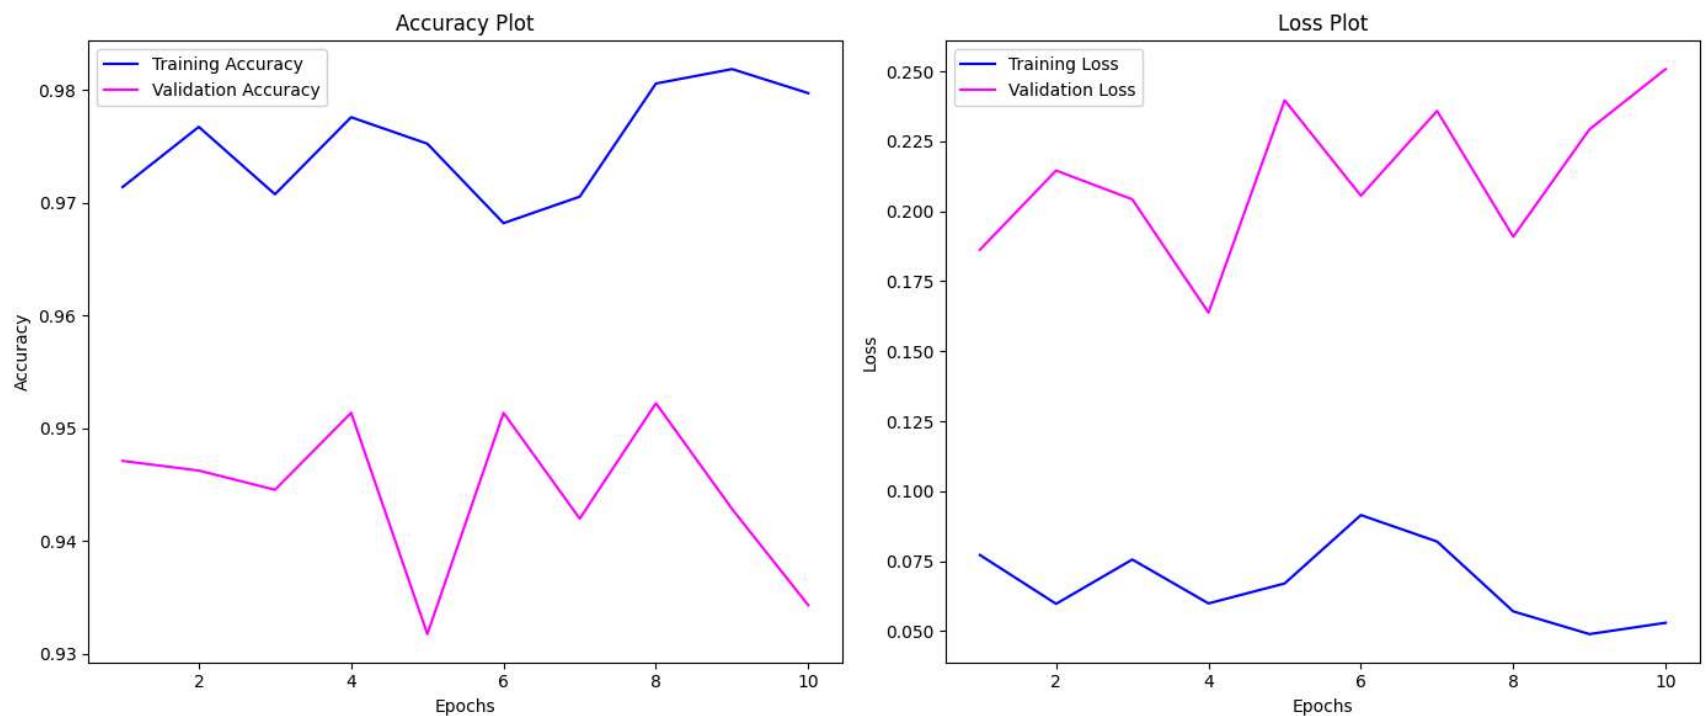
# Extract accuracy and error from history
accuracy = history.history['accuracy']
val_accuracy = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(accuracy) + 1)

# Draw the accuracy graph
plt.figure(figsize=(14, 6))

plt.subplot(1, 2, 1)
plt.plot(epochs, accuracy, color='blue', label='Training Accuracy')
plt.plot(epochs, val_accuracy, color='magenta', label='Validation Accuracy')
plt.title('Accuracy Plot')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Draw an error diagram
plt.subplot(1, 2, 2)
plt.plot(epochs, loss, color='blue', label='Training Loss')
plt.plot(epochs, val_loss, color='magenta', label='Validation Loss')
plt.title('Loss Plot')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```



```
In [11]: import numpy as np
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Predicting classes on test data
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

# Calculation of confusion matrix
cm = confusion_matrix(y_true, y_pred_classes)
tn, fp, fn, tp = cm.ravel()

# Calculation of sensitivity and specificity
sensitivity = tp / (tp + fn)
specificity = tn / (tn + fp)

print(f'Sensitivity: {sensitivity:.2f}')
print(f'Specificity: {specificity:.2f}')

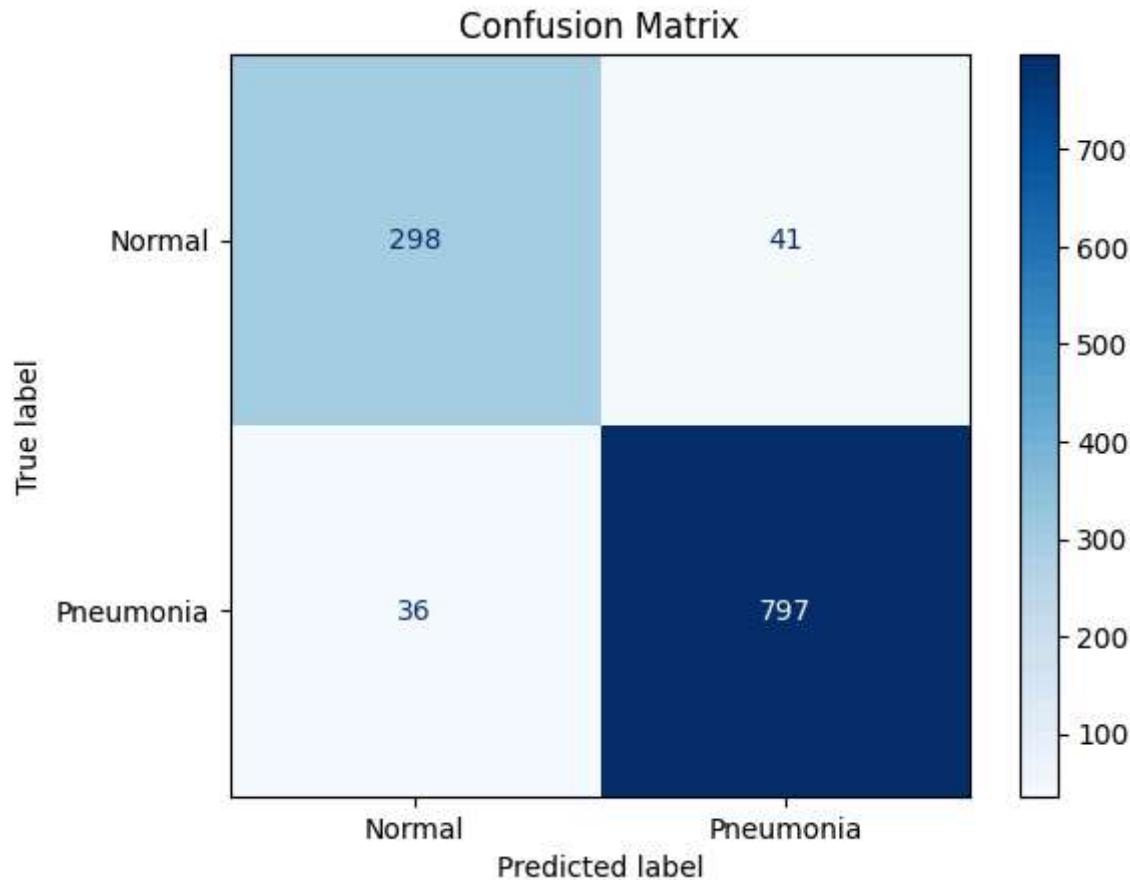
# Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Normal', 'Pneumonia'])
disp.plot(cmap='Blues')
plt.title('Confusion Matrix')
plt.show()

# Draw a sensitivity and specificity diagram
epochs = range(1, len(history.history['accuracy']) + 1)

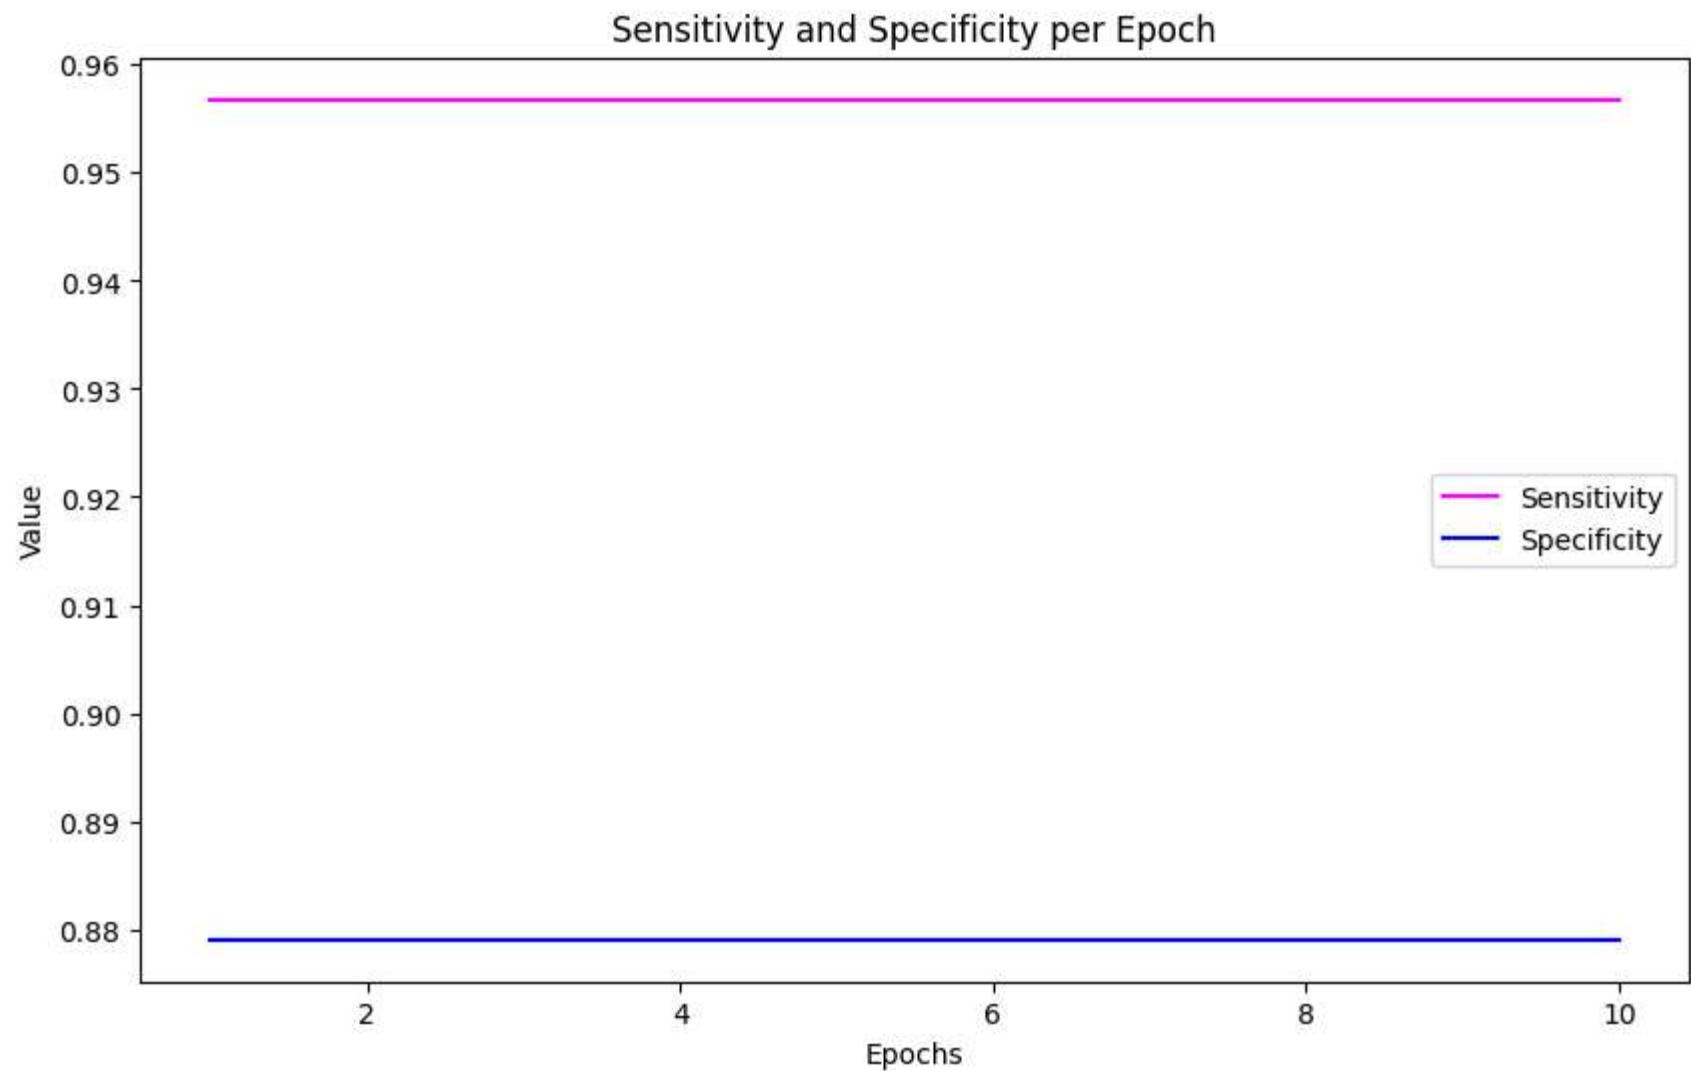
# Draw the diagram of sensitivity and specificity
plt.figure(figsize=(10, 6))
plt.plot(epochs, [sensitivity] * len(epochs), 'b', label='Sensitivity', color='magenta')
plt.plot(epochs, [specificity] * len(epochs), 'b', label='Specificity', color='blue')
plt.title('Sensitivity and Specificity per Epoch')
plt.xlabel('Epochs')
plt.ylabel('Value')
plt.legend()
plt.show()
```

37/37 ━━━━━━ 0s 4ms/step

Sensitivity: 0.96
Specificity: 0.88



```
<ipython-input-11-705c81530d1c>:32: UserWarning: color is redundantly defined by the 'color' keyword argument  
and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argument will take precedence.  
    plt.plot(epochs, [sensitivity] * len(epochs), 'b', label='Sensitivity', color='magenta')  
<ipython-input-11-705c81530d1c>:33: UserWarning: color is redundantly defined by the 'color' keyword argument  
and the fmt string "b" (-> color=(0.0, 0.0, 1.0, 1)). The keyword argument will take precedence.  
    plt.plot(epochs, [specificity] * len(epochs), 'b', label='Specificity', color='blue')
```




```
In [12]: import matplotlib.pyplot as plt
import cv2
import numpy as np
import random

# Randomly select an image from the dataset
index = random.randint(0, len(images) - 1) # 'images' is a list of images
sample_image = images[index]
sample_label = labels[index] # True label related to the image

# Reshape the sample image for prediction
sample_image_for_prediction = np.expand_dims(sample_image, axis=0) # Expand dimensions to match model input

# Step 1: Model prediction for the selected image
predicted_proba = model.predict(sample_image_for_prediction)
predicted_label = np.argmax(predicted_proba, axis=1)[0] # Get predicted class label

# Determine disease based on predicted label
if predicted_label == 0:
    predicted_disease = "Normal"
else:
    predicted_disease = "Pneumonia"

# Display the image and the predicted disease
plt.imshow(cv2.cvtColor(sample_image, cv2.COLOR_BGR2RGB))
plt.axis('off') # Disable image axes
plt.title(f'Predicted: {predicted_disease}', fontsize=20, color='blue')
plt.show()

# Step 2: Now, check the true label and compare
if sample_label == 0:
    actual_disease = "Normal"
else:
    actual_disease = "Pneumonia"

print(f"Actual label: {actual_disease}")

# Step 3: Check if prediction is correct
if predicted_label == sample_label:
    print("Prediction is correct.")
else:
```

```
print("Prediction is incorrect.")
```

1/1 **1s** 550ms/step

Predicted: Pneumonia



Actual label: Pneumonia

Prediction is correct.

```
In [13]: # Increase image contrast
def increase_contrast(image):
    gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    equalized_image = cv2.equalizeHist(gray_image)
    return equalized_image

# Select an image
sample_image = images[index]
contrast_image = increase_contrast(sample_image)

# Display original image and image with high contrast
plt.subplot(1, 2, 1)
plt.imshow(cv2.cvtColor(sample_image, cv2.COLOR_BGR2RGB))
plt.title('Original Image')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(contrast_image, cmap='gray')
plt.title('Increased Contrast')
plt.axis('off')

plt.show()
```

Original Image



Increased Contrast



```
In [14]: # Reduce image noise
def reduce_noise(image):
    blurred_image = cv2.GaussianBlur(image, (5, 5), 0)
    return blurred_image

# noise reduction
noise_reduced_image = reduce_noise(sample_image)

# Display the original image and the image with noise reduction
plt.subplot(1, 2, 1)
plt.imshow(cv2.cvtColor(sample_image, cv2.COLOR_BGR2RGB))
plt.title('Original Image')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(cv2.cvtColor(noise_reduced_image, cv2.COLOR_BGR2RGB))
plt.title('Noise Reduced')
plt.axis('off')

plt.show()
```

Original Image



Noise Reduced




```
In [15]: import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

# 1. Upload images with pre-processing
def load_images(data_path):
    images = []
    labels = []

    for label in ['NORMAL', 'PNEUMONIA']:
        label_path = os.path.join(data_path, label)
        for img_file in os.listdir(label_path):
            img_path = os.path.join(label_path, img_file)
            img = cv2.imread(img_path)

            # Convert to gray image
            img_gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

            # noise reduction
            img_denoised = cv2.GaussianBlur(img_gray, (5, 5), 0)

            # Increase contrast
            img_eq = cv2.equalizeHist(img_denoised)

            # Resize to 128x128
            img_resized = cv2.resize(img_eq, (128, 128))
            images.append(img_resized)
            labels.append(int(label == 'PNEUMONIA'))

    return np.array(images), np.array(labels)

# Upload images
data_path = '/content/chest-xray-pneumonia/chest_xray/train'
images, labels = load_images(data_path)

# 2. Dividing data into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(images, labels, test_size=0.2, random_state=42)
```

```
# 3. Data augmentation
datagen = ImageDataGenerator(
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True)

# 4. Create a model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 1)),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(2, activation='softmax')
])

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# 5. Model training
history = model.fit(datagen.flow(X_train[..., np.newaxis], y_train, batch_size=32),
                     validation_data=(X_val[..., np.newaxis], y_val),
                     epochs=10)

# 6. Model evaluation
test_loss, test_accuracy = model.evaluate(X_val[..., np.newaxis], y_val)
print(f'Test accuracy: {test_accuracy:.2f}')
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Epoch 1/10

/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self._warn_if_super_not_called()

```
131/131 13s 70ms/step - accuracy: 0.7097 - loss: 53.2976 - val_accuracy: 0.8621 - val_loss: 0.3833
Epoch 2/10
131/131 8s 56ms/step - accuracy: 0.8798 - loss: 0.3126 - val_accuracy: 0.8946 - val_loss: 0.2969
Epoch 3/10
131/131 7s 49ms/step - accuracy: 0.8796 - loss: 0.3061 - val_accuracy: 0.9109 - val_loss: 0.2462
Epoch 4/10
131/131 7s 51ms/step - accuracy: 0.8975 - loss: 0.2567 - val_accuracy: 0.9071 - val_loss: 0.2360
Epoch 5/10
131/131 7s 54ms/step - accuracy: 0.9108 - loss: 0.2287 - val_accuracy: 0.9080 - val_loss: 0.2894
Epoch 6/10
131/131 7s 46ms/step - accuracy: 0.9166 - loss: 0.2208 - val_accuracy: 0.9195 - val_loss: 0.2195
Epoch 7/10
131/131 8s 59ms/step - accuracy: 0.9226 - loss: 0.1984 - val_accuracy: 0.9186 - val_loss: 0.2075
Epoch 8/10
131/131 6s 46ms/step - accuracy: 0.9213 - loss: 0.2040 - val_accuracy: 0.9310 - val_loss: 0.1953
Epoch 9/10
131/131 8s 56ms/step - accuracy: 0.9249 - loss: 0.1934 - val_accuracy: 0.9387 - val_loss: 0.1755
Epoch 10/10
131/131 9s 47ms/step - accuracy: 0.9296 - loss: 0.1752 - val_accuracy: 0.9033 - val_loss: 0.2520
33/33 0s 4ms/step - accuracy: 0.9009 - loss: 0.2465
Test accuracy: 0.90
```

In [16]: # Random forest

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split

# Dividing data into training and testing sets (pre-processing as before)
X_train, X_test, y_train, y_test = train_test_split(images.reshape(len(images), -1), labels, test_size=0.2, random_state=42)

# Construction of Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

# Model training
rf_model.fit(X_train, y_train)

# Prediction on test data
y_pred_rf = rf_model.predict(X_test)

# Calculate accuracy
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print(f'Random Forest Validation Accuracy: {accuracy_rf * 100:.2f}%')

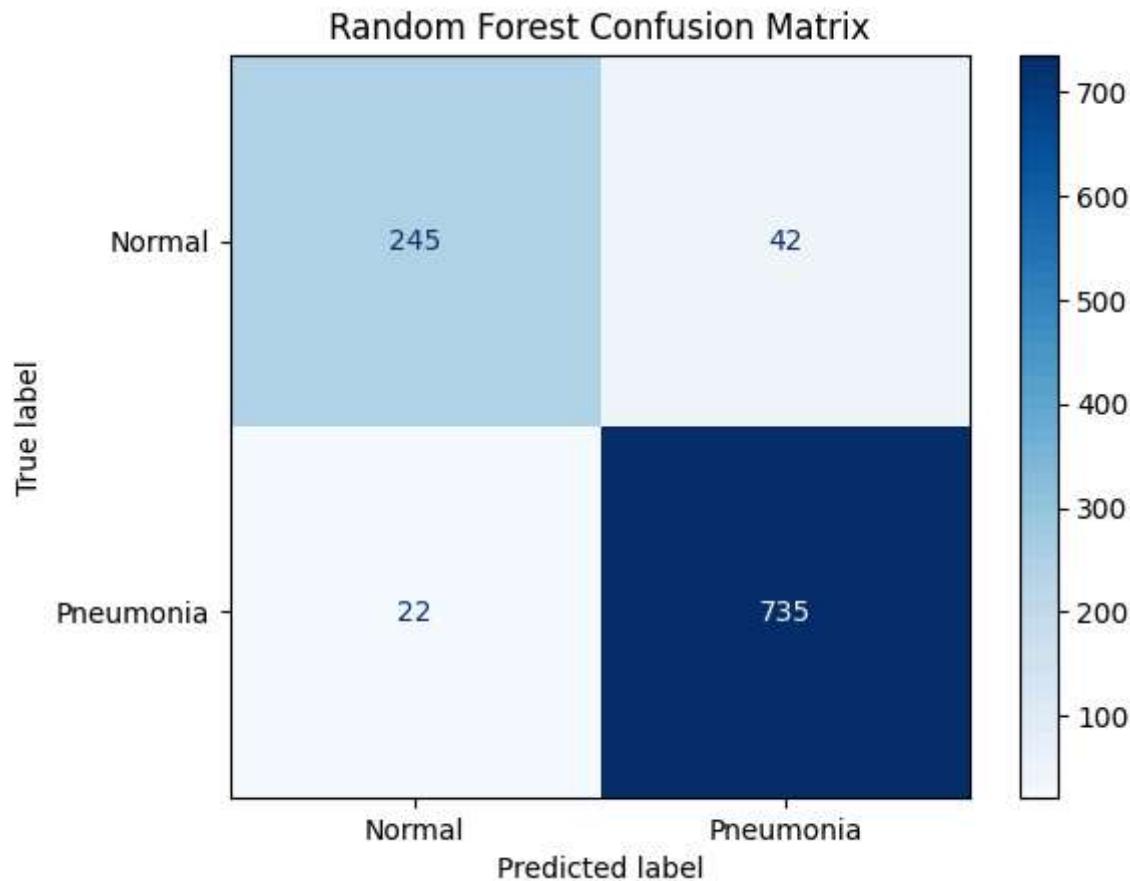
# Calculation of confusion matrix
cm_rf = confusion_matrix(y_test, y_pred_rf)
tn_rf, fp_rf, fn_rf, tp_rf = cm_rf.ravel()

# Calculation of sensitivity and specificity
sensitivity_rf = tp_rf / (tp_rf + fn_rf)
specificity_rf = tn_rf / (tn_rf + fp_rf)

print(f'Random Forest Sensitivity: {sensitivity_rf:.2f}')
print(f'Random Forest Specificity: {specificity_rf:.2f}')

# Display the confusion matrix
disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf, display_labels=['Normal', 'Pneumonia'])
disp_rf.plot(cmap='Blues')
plt.title('Random Forest Confusion Matrix')
plt.show()
```

Random Forest Validation Accuracy: 93.87%
Random Forest Sensitivity: 0.97
Random Forest Specificity: 0.85



In [17]: # Precision vs Recall (Random Forest)

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_recall_curve, classification_report
import matplotlib.pyplot as plt
import numpy as np

# Random Forest model training
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Prediction of tags
y_pred_rf = rf_model.predict(X_test)

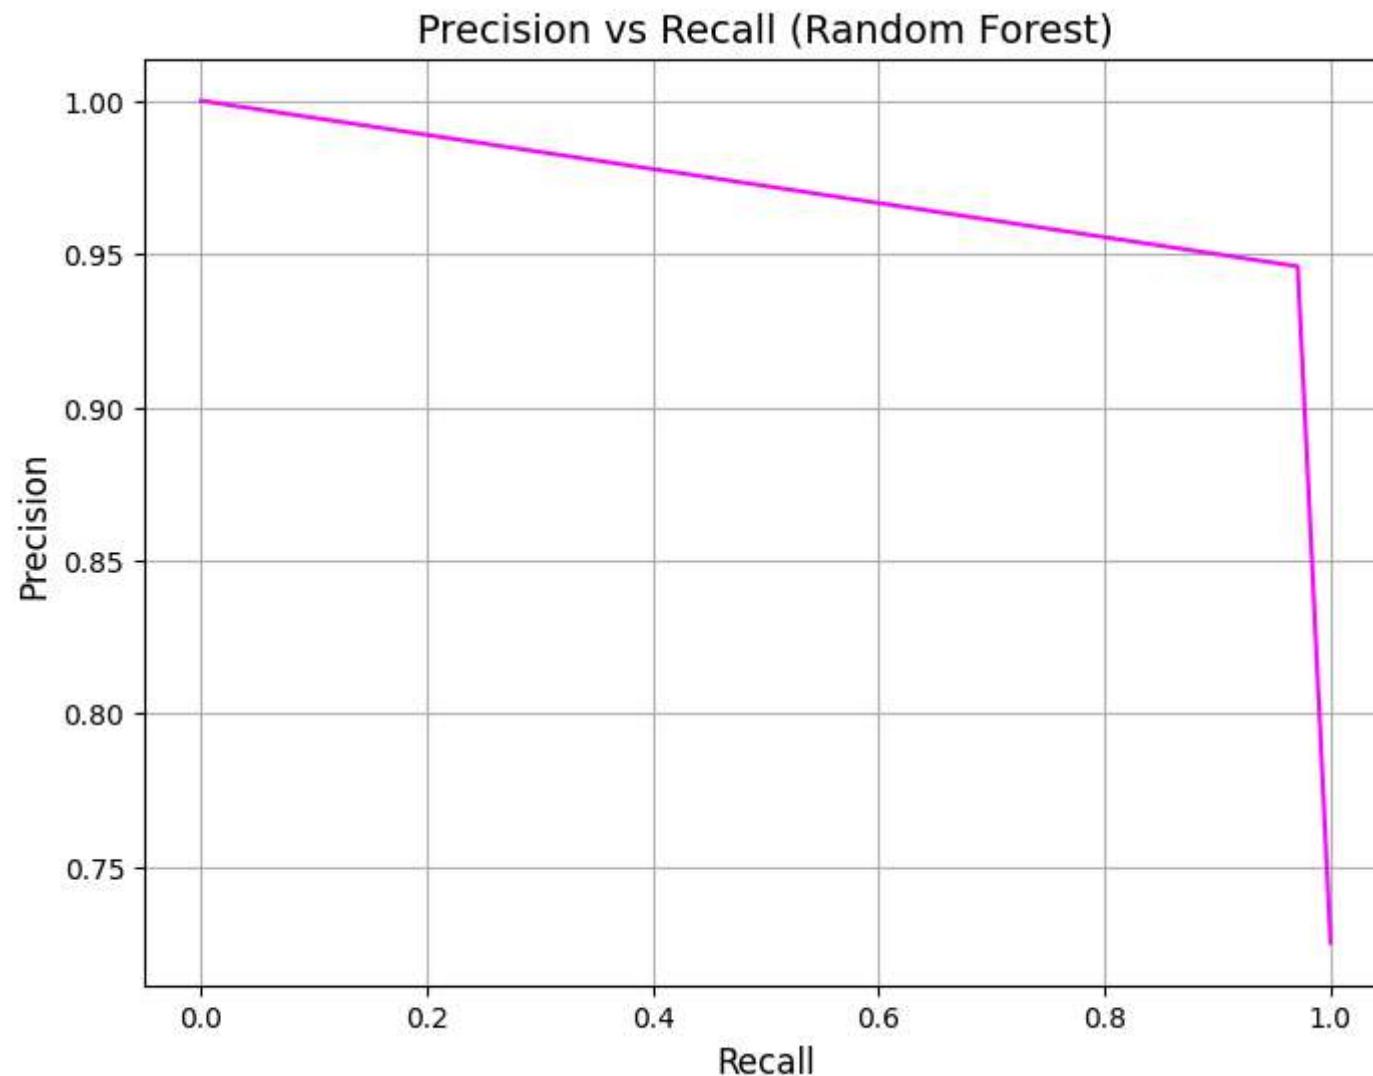
# Calculation of precision and recall
precision_rf, recall_rf, _ = precision_recall_curve(y_test, y_pred_rf)

# Full report
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))

# Draw a Precision vs Recall graph
plt.figure(figsize=(8, 6))
plt.plot(recall_rf, precision_rf, color='magenta')
plt.title('Precision vs Recall (Random Forest)', fontsize=14)
plt.xlabel('Recall', fontsize=12)
plt.ylabel('Precision', fontsize=12)
plt.grid(True)
plt.show()
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.92	0.85	0.88	287
1	0.95	0.97	0.96	757
accuracy			0.94	1044
macro avg	0.93	0.91	0.92	1044
weighted avg	0.94	0.94	0.94	1044



In [18]:

```
# XGBOOST
!pip install xgboost
import xgboost as xgb
from sklearn.metrics import accuracy_score, confusion_matrix

# Building the XGBoost model
xgb_model = xgb.XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')

# Model training
xgb_model.fit(X_train, y_train)

# Prediction on test data
y_pred_xgb = xgb_model.predict(X_test)

# Calculate accuracy
accuracy_xgb = accuracy_score(y_test, y_pred_xgb)
print(f'XGBoost Validation Accuracy: {accuracy_xgb * 100:.2f}%')

# Calculation of confusion matrix
cm_xgb = confusion_matrix(y_test, y_pred_xgb)
tn_xgb, fp_xgb, fn_xgb, tp_xgb = cm_xgb.ravel()

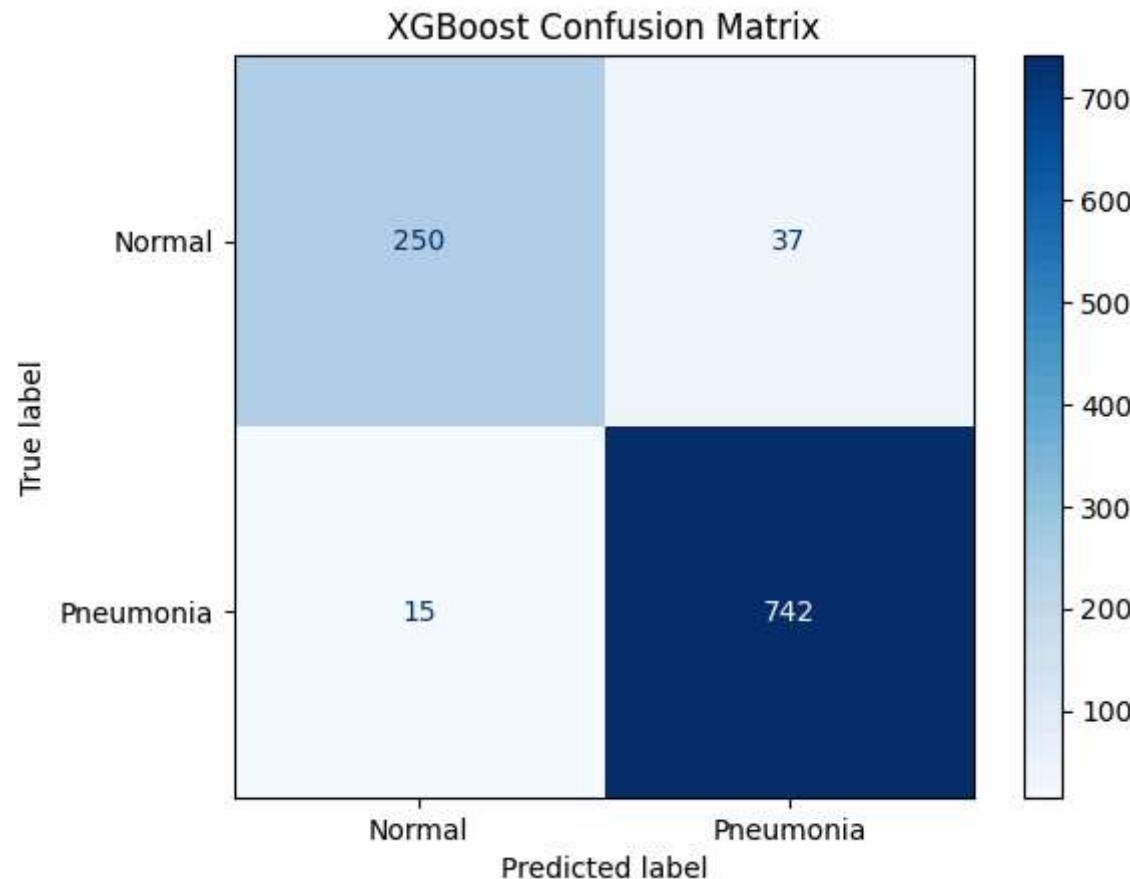
# Calculation of sensitivity and specificity
sensitivity_xgb = tp_xgb / (tp_xgb + fn_xgb)
specificity_xgb = tn_xgb / (tn_xgb + fp_xgb)

print(f'XGBoost Sensitivity: {sensitivity_xgb:.2f}')
print(f'XGBoost Specificity: {specificity_xgb:.2f}')

# Display the confusion matrix
disp_xgb = ConfusionMatrixDisplay(confusion_matrix=cm_xgb, display_labels=['Normal', 'Pneumonia'])
disp_xgb.plot(cmap='Blues')
plt.title('XGBoost Confusion Matrix')
plt.show()
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-packages (2.1.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.26.4)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.10/dist-packages (from xgboost) (2.2
3.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.13.1)
```

```
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning: [17:18:46] WARNING: /workspace/src/learner.cc:740:  
Parameters: { "use_label_encoder" } are not used.  
  
warnings.warn(smsg, UserWarning)  
  
XGBoost Validation Accuracy: 95.02%  
XGBoost Sensitivity: 0.98  
XGBoost Specificity: 0.87
```



In [19]: # Precision vs Recall (XGBOOST)

```
from xgboost import XGBClassifier
from sklearn.metrics import precision_recall_curve, classification_report
import matplotlib.pyplot as plt
import numpy as np

# XGBoost model training
xgb_model = XGBClassifier(random_state=42)
xgb_model.fit(X_train, y_train)

# Prediction of tags
y_pred_xgb = xgb_model.predict(X_test)

# Calculation of precision and recall
precision_xgb, recall_xgb, _ = precision_recall_curve(y_test, y_pred_xgb)

# Full report
print("XGBoost Classification Report:")
print(classification_report(y_test, y_pred_xgb))

# Draw a Precision vs Recall graph
plt.figure(figsize=(8, 6))
plt.plot(recall_xgb, precision_xgb, color='blue')
plt.title('Precision vs Recall (XGBoost)', fontsize=14)
plt.xlabel('Recall', fontsize=12)
plt.ylabel('Precision', fontsize=12)
plt.grid(True)
plt.show()
```

XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.94	0.87	0.91	287
1	0.95	0.98	0.97	757
accuracy			0.95	1044
macro avg	0.95	0.93	0.94	1044
weighted avg	0.95	0.95	0.95	1044

